

Report on Data Centers

Yongsen MA

I. MOTIVATION

With the increasing demand of traffic loads and user requirements, it is a great challenge to make data centers agile and energy-efficient. A basic solution is to allow dynamic resource allocation based on flexible data center networks. The recent development of 60GHz wireless technology opens a door for the deployment of flexible data centers. It provides a good choice of adding capacity to data centers with under-provisioning capacity design. However, this will lead to new problems such as topology control and capacity scheduling, apart from the wireless propagation and link adaption in wireless networks. The PHY and MAC standardization of 60GHz networks is still ongoing (WirelessHD and IEEE 802.11ad/WiGig), its application in data centers should be further explored for the specific feature of topology and services.




Fig. 1. Capacity provisioning based on traffic loads

Generally, the capacity of data centers is designed for peak loads, as shown in Figure 1a, but the resources will be wasted during non-peak times [Armbrust et al., 2010]. On the other hand, if the capacity is designed in under-provisioning case as shown in Figure 1b, the peak requirements can not be satisfied, leading to potential revenue sacrifice or over-occupied errors. To address these problems, we can employ flexible capacity and dynamic demand response, i.e., 60GHz wireless linking to add capacity for peak loads and sleeping mechanism to improve performance/cost efficiency, as shown in Figure 1c.

The wireless-flexible data centers are generally composed of:

- 1) Topology and Addressing:
 - a) Conflict Graph: Propagation table (wireless), capacity matrix (wired/wireless) and traffic loads table (wired/wireless)
 - b) Topology Graph: Physical-specific ID Addresses (PAs), Logical-specific IP Addresses (LAs) and Application-specific IP Addresses (AAs)
 - i) PAs: providing distance information for 3D Beamforming
 - ii) LAs: providing topology information for Capacity Adaption
 - iii) AAs: providing traffic information for Demand Response
- 2) Routing and Scheduling:
 - a) 3D Beamforming: adding additional capacity to neighbor ToRs according to PAs
 - b) Demand Response: traffic estimation according to AAs
 - c) Capacity Adaption: capacity scheduling according to LAs

II. METHODOLOGY

The problems of topology control, addressing, routing, and scheduling also exist in traditional data centers, but it brings new challenges when wireless links are introduced. For instance, the topology graph is changing due to wireless links and on-demand response. Also, the capacity matrix and conflict graph of wireless links are closely related to the relative locations of servers, e.g., 6G for neighbor servers and 1/2G for non-neighbor servers. Furthermore, the measurement and mapping of topology graph, conflict graph and traffic loads is challenging for large scale data centers.

A. Measurement

- Propagation table
- Capacity matrix
- Traffic loads

Topology changes:

- Errors: nodes, links, miswiring
- wireless linking
- Sleeping on-off

Mapping:

- PA to LA (Beamforming)
- AA to LA (Traffic estimation)

Connections:

- wired 10G
- wireless direct 6G
- wireless indirect 1/2G

wireless data centers: making data centers flexible, i.e., using topology control to make resource allocation responding to traffic patterns and requirements.

- graph mapping: topology changes due to errors or wireless links.
- demand response: topology and capacity (wired and wireless) are stable for certain graphs, so the problem is demand estimation.

So how to allocate wireless links (direct links between neighbor nodes or 3D beamforming for remote nodes), to deal with the problem of demand response when topology graphs are changing. Topology changes may occur due to errors, sleeping mechanism or wireless links.

reasons for compress sensing:

- 1) topology graphs are **large**: the requirement for compress sensing to reduce mapping overhead
- 2) wireless propagation and traffic demand measurement is challenging: requirements
- 3) topology graphs are **sparse**: the prior condition for compress sensing
- 4) topology graphs have **local changes** (errors, on-off or wireless) in real-time operating: conditions

B. Scheduling

III. LITERATURE REVIEW

Data Center Networks:

- Static-electrical:
 - Server-centric:
 - * DCell [Guo et al., 2008]
 - * BCube [Guo et al., 2009]
 - Multi-level tree:
 - * Fat-tree [Al-Fares et al., 2008]
 - * VL2 [Greenberg et al., 2009]
- Flexible:
 - Optical:
 - * c-Through [Wang et al., 2010]
 - * Helios [Farrington et al., 2010]
 - * OSA [Chen et al., 2012]
 - Wireless:
 - * Wireless flyways [Halperin et al., 2011]
 - * 3D beamforming [Zhou et al., 2012]

"Generic and Automatic Address Configuration for Data Center Networks", SIGCOMM 2010

Basic Procedures: [Chen et al., 2010]

- 1) O2 Mapping
 - a) Candidate selection via SPLD: **select** candidate with the same SPLD.
 - b) Candidate filtering via orbit: **skip** candidate with the same orbit, then *Decomposition()*.
 - c) Selective splitting *Refinement**(): **split** cells that really connect to the including cell.
- 2) Malfunction Detection
 - a) Anchor pair selection:
 - b) Malfunction detection:

Problems:

- 1) Initial selection of vertex $\nu \in \pi_p^i$
- 2) Whether it can be resolved by Compress Sensing?
 - a) the topology graphs are sparse.
 - b) only certain parts are changing in real-time operating (considering certain servers can be turned down for energy-efficiency and demand response).

"OSA: An Optical Switching Architecture for Data Center Networks with Unprecedented Flexibility", NSDI 2012

Architecture

"A Scalable, Commodity Data Center Network Architecture", SIGCOMM 2008

Topology: k pods, two layers of $k/2$ switches, each k -port switch connected to $k/2$ hosts, each remaining $k/2$ ports connected to $k/2$ of the k ports in the aggregation layer of the hierarchy. $\text{Max } 48 \cdot \frac{48}{2} \cdot \frac{48}{2} = 27648$ hosts for 48-port GigE switches.

Addressing: Pod: 10.pod.switch.1, core: 10.k.j.i, host: 10.pod.switch.ID

"DCell: A Scalable and Fault-Tolerant Network Structure for Data Centers", SIGCOMM 2008

Large-scale, server-centric data centers

Topology: n servers connected to a mini-switch within $DCell_0$, then $DCell_k$ is recursively generated by $t_{k-1} + 1$ $DCell_{k-1}$ s. The number of servers scales doubly exponentially as the node degree increases.

Addressing: a server in $DCell_k$ is denoted by $[a_k, uid_{k-1}]$, where a_k is the $DCell_k$ this server belongs to and uid_{k-1} is the unique ID of the server inside this $DCell_{k-1}$.

Routing: find the connection (n_1, n_2) between $DCells$, then find the paths of src to n_1 and n_2 to dst . When failures occur, choosing *proxy* to make local-reroute or jump-up (rack failure).

Problems:

- the bottleneck link locates in the low-level links and traffics are not balanced
- increment expansion requires additional ports which brings extra overhead and complexity

"BCube: A High Performance, Server-centric Network Architecture for Modular Data Centers", SIGCOMM 2009

Server-centric modular data centers

Topology: $BCube_0$ is n servers connecting to an n -port switch. $BCube_k$ is constructed from n $BCube_{k-1}$ s and n^k n -port switches.

Addressing: servers: $a_k a_{k-1} \cdots a_0$ ($a_i \in [0, n-1], i \in [0, k]$); switches: $\langle l, s_{k-1}, s_{k-2} \cdots s_0 \rangle$, where $l(0 \leq l \leq k)$ is the level of the switch.

Routing:

"VL2: A Scalable and Flexible Data Center Network", SIGCOMM 2009

1. Background:

- Limited server-to-server capacity: over-subscription ratio increases rapidly
- Fragmentation of resources: high turnaround time for reconfiguration
- Poor reliability and utilization: multiple paths waste at most 50% of maximum utilization

2. Measurement:

- Flow: VLB will perform well on this traffic
 - Distribution of flow sizes: majority of flows are small
 - Number of concurrent flows: 10 concurrent flows at more than 50% of the time, 80 concurrent flows at least 5% of the time
- Traffic: it is unlikely that other routing strategies will outperform VLB
 - Poor summarizability of traffic patterns
 - Instability of traffic patterns
- Failure

3. Solution: Virtual Layer 2 Networking

- Topology: D_I Int. Switches and D_A Agg. Switches
- Addressing: Location-specific IP Addresses and Application-specific IP Addresses
- Directory: lookups and updates of AA-to-LA mapping, reactive cache update

Problems:

- arrive intervals.
- application types.

Wireless

"Mirror Mirror on the Ceiling: Flexible Wireless Links for Data Centers", SIGCOMM 2012

Wireless Data Center is a new concept which should be explored further. Many issues such as interference, secluding, security, etc. should be standardized. On the other hand, the issues on cost,

performance, energy-efficiency, reliability, etc. should be taken into consideration compared with electrical or optical data centers.

- 1) electrical
- 2) optical
- 3) wireless
 - a) Link Blockage
 - b) Radio Interference

Max concurrent links: Link conflicts (SINR); Greedy scheduling (graph coloring); Assigning radios.

Problems:

- Apart from the concurrent links, the efficient throughput should be explored. For instance, although the concurrent links are more with larger ceiling height h , it can decrease the throughput according to the curves of RSS (or Data Rate) vs. distance as shown in Figure 5.

"Augmenting Data Center Networks with Multi-gigabit Wireless Links", SIGCOMM 2011

"The base wired network is provisioned for the **average case** and can be oversubscribed. Each ToR switch is equipped with one or more 60GHz wireless devices." "A central controller monitors DC **traffic patterns**, and switches the beams of the wireless devices to set up flyways between ToR switches that provide **added bandwidth** as needed." So it is ideal for **flexible and energy-efficient** data centers.

Problems:

- Conflict graph: for N racks and K antenna orientations, the input table is very **large** with the size of $(NK)^2$. On the other hand, the propagation conditions are similar, i.e., the table is **sparse**.

Resource Allocation (to be done)

"FairCloud: Sharing The Network In Cloud Computing", SIGCOMM 2012

"NetPilot: Automating Datacenter Network Failure Mitigation", SIGCOMM 2012

"Towards Predictable Datacenter Networks", SIGCOMM 2011

REFERENCES

- [Al-Fares et al., 2008] Al-Fares, M., Loukissas, A., and Vahdat, A. (2008). A scalable, commodity data center network architecture. *ACM SIGCOMM '08*.
- [Armbrust et al., 2010] Armbrust, M., Fox, A., Griffith, R., Joseph, A. D., Katz, R., Konwinski, A., Lee, G., Patterson, D., Rabkin, A., Stoica, I., and Zaharia, M. (2010). A view of cloud computing. *Commun. ACM*, 53(4):50–58.
- [Chen et al., 2010] Chen, K., Guo, C., Wu, H., Yuan, J., Feng, Z., Chen, Y., Lu, S., and Wu, W. (2010). Generic and automatic address configuration for data center networks. *ACM SIGCOMM '10*.
- [Chen et al., 2012] Chen, K., Singla, A., Singh, A., Ramachandran, K., Xu, L., Zhang, Y., Wen, X., and Chen, Y. (2012). OSA: An optical switching architecture for data center networks with unprecedented flexibility. *USENIX NSDI '12*.
- [Farrington et al., 2010] Farrington, N., Porter, G., Radhakrishnan, S., Bazzaz, H. H., Subramanya, V., Fainman, Y., Papen, G., and Vahdat, A. (2010). Helios: a hybrid electrical/optical switch architecture for modular data centers. In *Proceedings of the ACM SIGCOMM 2010 conference*, ACM SIGCOMM '10.
- [Greenberg et al., 2009] Greenberg, A., Hamilton, J. R., Jain, N., Kandula, S., Kim, C., Lahiri, P., Maltz, D. A., Patel, P., and Sengupta, S. (2009). VL2: a scalable and flexible data center network. *ACM SIGCOMM '09*.
- [Guo et al., 2009] Guo, C., Lu, G., Li, D., Wu, H., Zhang, X., Shi, Y., Tian, C., Zhang, Y., and Lu, S. (2009). BCube: a high performance, server-centric network architecture for modular data centers. *ACM SIGCOMM '09*.
- [Guo et al., 2008] Guo, C., Wu, H., Tan, K., Shi, L., Zhang, Y., and Lu, S. (2008). DCell: a scalable and fault-tolerant network structure for data centers. *ACM SIGCOMM '08*.
- [Halperin et al., 2011] Halperin, D., Kandula, S., Padhye, J., Bahl, P., and Wetherall, D. (2011). Augmenting data center networks with multi-gigabit wireless links. *ACM SIGCOMM '11*.
- [Wang et al., 2010] Wang, G., Andersen, D. G., Kaminsky, M., Papagiannaki, K., Ng, T. E., Kozuch, M., and Ryan, M. (2010). c-Through: part-time optics in data centers. In *Proceedings of the ACM SIGCOMM 2010 conference*, ACM SIGCOMM '10.

[Zhou et al., 2012] Zhou, X., Zhang, Z., Zhu, Y., Li, Y., Kumar, S., Vahdat, A., Zhao, B. Y., and Zheng, H. (2012). Mirror mirror on the ceiling: flexible wireless links for data centers. ACM SIGCOMM '12.