Report on Data Centers

Yongsen MA

I. LITERATURE REVIEW

According to the architecture and topology in these papers, Data Center Networks can be divided into:

1) Static-electrical:
   a) Multi-layer tree:
      i) Fat-tree [1]
      ii) VL2 [2]
   b) Server-centric:
      i) DCell [3]
      ii) BCube [4]
   c) Random graph:
      i) Jellyfish [5]
2) Flexible:
   a) Optical:
      i) c-Through [9]
      ii) Helios [10]
      iii) OSA [11]
   b) Wireless:
      i) Wireless flyways [12]
      ii) 3D beamforming [13]
      iii) Cylindrical racks [14]

These papers mainly discussed the issues of:

- Topology [1]–[5]
- Addressing [1], [2]
- Mapping [2], [15]
- Routing [3], [4]
- Scheduling [1], [12]
- Failure handling [15]–[17]
- Multi-tenant [6]–[8]

A. Static-electrical Data Centers

The traditional data centers are generally designed in static connection by electrical switches. For the issues of topology, routing and scheduling in these data centers, there are numerous researches in this field including Fat-tree [1], VL2 [2], DCell [3], BCube [4], Jellyfish [5], etc. These work can be classified into multi-layer tree, server-centric and random graph according to the differences in topology graph, which also brings different advantages and disadvantages.

1) Multi-layer Tree: The topology of fat-tree is multi-rooted tree which can provide full aggregate bandwidth with low cost and resolved the bottleneck in inter-node communication. But the scale of fat-tree is limited that 48-port switches can only support $48 \cdot \frac{48}{2} \cdot \frac{48}{2} = 27648$ hosts at most.

The conventional data center architecture has limited server-to-server capacity that the over-subscription ratio increases rapidly in upper layer switches. Also, servers in multi-layer tree are location related that nearby servers are responsible for demand peaks or failures resulting in resources wasting. According to the measurement and analysis of traffic patterns, VL2 adopts Virtual Layer 2 Networking to realize the dynamic resource allocation across large server pools. It can assign any server to any service which can address the fluctuating demands of individual services.

The multi-rooted tree topology has poor fault-tolerant features which make it hard to detect and maintain failures. Furthermore, multi-rooted tree can not support one-to-all and all-to-all communications effectively.

2) Server-centric: For server-centric data centers, all-to-all communication can be realized and the network can be incrementally expanded. DCell is a recursively defined structure which scales doubly exponentially as the node degree increases (ideal for large scale data centers). Furthermore, DCell is fault tolerant and can provide higher network capacity. But the bottleneck link locates in the low-level links and traffics are not balanced. Moreover, increment expansion requires additional ports which brings extra overhead and complexity.

BCube provides a similar server-centric solution for modular data centers, which has no performance bottlenecks and provides much higher network capacity. However, to connect certain number of servers, it needs more switches (k times) and wiring is more complex which will increase the cabling cost. For BCube, the aggregate bottleneck throughput decreases gracefully under both server and switch failures, which means poor performance when incomplete units exit.

In the multi-path routing and speedup forwarding of one-to-all communications, it needs **strict time synchronization**, and more importantly, brings the trade-off of reliability and capacity just as MIMO in wireless networks.

3) Random Graph: The architectures listed above have rigid structure that interferes with incremental expansion. Jellyfish adopts random graph topology to address this problem, and also provides great flexibility in oversubscription design. But it brings challenging complexity to routing, wiring and maintaining.

B. Flexible Data Centers

1) Optical: Optical circuit switching technology has large bandwidth advantage over conventional packet switching, which provides suitable solution for data centers. But optical networking has **slow switching** that will increase the reconfiguration time, which remains as the main challenge for optical data centers. To take full advantage of high capacity of optical switching while ensure low latency, c-Through proposed a hybrid packet and circuit switched data center network architecture. Along with the hybrid architecture, it also makes demand estimation and traffic demultiplexing in end hosts. But it has high **configuration complexity** in separating the networks into packet/circuit-switch and the preconnection socket buffer. The socket buffers at end hosts also **increase the overhead** of traffic measurement and software programming. c-Through requires each host run a management daemon which further increases the system overhead. Finally, the authors use an Ethernet switch to emulate the optical switch and artificially restrict the communication and configuration, it will reduce the accuracy of evaluating the influence of optical's high latency.

Unlike c-Through, Helios implements traffic estimation and demultiplexing on switches and requires no modifications to end hosts. The topology of Helios is multi-rooted tree in which core switches can

be either electrical packet switches or optical circuit switches. The optical parts **add flexibility** to data centers which can get optimal trade-off of cost, power consumption, complexity and performance for a given set of workloads. In the evaluation of Helios, the authors disable the function of debouncing and Electronic Dispersion Compersation (EDC) to improve PHY throughput. But this will perhaps reduce the upper layer **Goodput** due to link insertion/removal and signal noise.

The hybrid structure in OSA is an abstract model that captures key aspects of both c-Through and Helios. In addition, OSA can connect one ToR to multiple ToRs simultaneously and multi-hop connection exists between any pair of remote ToRs.

2) Wireless: Another solution for flexible data centers is 60GHz wireless links, but wireless data center is a new concept which should be explored further. Many issues such as interference, secluding, security, etc. should be standardized. On the other hand, the issues on cost, performance, energy-efficiency, reliability, etc. should be taken into consideration compared with electrical and optical data centers.

Apart from the physical topology, the conflict graph should be measured when wireless links are introduced. For N racks and K antenna orientations, the input table of conflict graph is very **large** with the size of $(NK)^2$. On the other hand, the propagation conditions are similar, which means the conflict table is **sparse**. In addition to wireless propagation, upper layer link state should also be designed specifically for data center networks. Apart from the concurrent links characterized in [13], the **efficient throughput** should be explored. For instance, although the concurrent links are more with larger ceiling height h, it can decrease the throughput according to the curves of RSS (or Data Rate) vs. distance.

C. Other issues

1) Multi-tenant: When data centers are shared in the cloud, i.e., multi-tenant data centers, the resource allocation is challenging due to the variability of tenants, application types and traffic patterns. In addition to the optimization within data center networking, the cloud providers should also conduct scheduling between different tenants. Oktopus [7] adopts virtual network abstractions to make prediction and allocation. NetLord [6] also provides tenants with flexible network abstractions by fully and efficiently virtualization of he address at both layer 2 and 3. It requires only static configuration and local information, but it will increase the **reconfiguration complexity** when there is topology change which will occur frequently in flexible data centers. FairCloud [8] exposes the fundamental trade-offs in multi-tenant data centers and develops a set of resource allocation policies to best navigate these trade-offs. It allocates the bandwidth according to the number of VMs rather than the number of flows, sources, or source destination pairs. But the proportional allocation of link sharing is difficult in practical implementation.

2) Failure Handling: For large scale data centers, there is a fundamental problem of failure handling including detection, diagnosis and repair. NetPilot [17] presents an automated system to quickly mitigate failures before the diagnose and repair process. It takes advantage of **multi-level redundancy** to mitigate most types of failures by simple actions such as deactivation or restart.

DAC [15] presents fast device-to-logical mapping with low time complexity by leveraging the properties of topology graph, and it can also help to efficiently detect errors such as link failures and miswirings. Since DAC is generally designed for different data center structures, it can be further developed for certain topologies. But DAC does not consider capacity matrix including differences in up- and down-link (directed graph), and in connections (weighted graph).

Problems:

1) Initial selection of vertex $\nu \in \pi_p^i$
2) Whether it can be resolved by Compress Sensing?
   a) the topology graphs are sparse.

b) only certain parts are changing in real-time operating (considering certain servers can be turned down for energy-efficiency and demand response).

II. RESEARCH ISSUES

A. Topology Control

1) topology graph is large and sparse
2) topology graph is changing due to errors, multiplexing and demand response
3) capacity differences in up- and down-link (directional graph)
4) capacity differences in connections (weighted graph)

It is challenging to measure such a large scale topology graph which usually requires all-to-all communications. Since both traffic matrix [2], [12] and topology graph [15] are sparse, we can utilize Compress Sensing to reduce overhead, just as the famous example that using only 3 out of 10 words can reconstruct the whole account password. In this way, we can use small number of samples to get the topology so that all-to-all communications are not necessary.

The requirements and reasons for Compress Sensing are:

1) topology graphs are **large**, and so comes the requirement to reduce measuring and mapping overhead
2) wireless propagation and traffic demand have **dynamic** furthers, so the measurement is challenging
3) topology graphs are **sparse**, which is the prior condition for compress sensing
4) topology graphs only have **local changes** (errors, on-off or wireless) in real-time operating

If the adjacency matrix of a given topology graph $G = (V, E)$ is $A^{N \times N}$, it can be represented by N vertex A_i. Then the estimated matrix $\hat{y} \in \mathbf{R}^n$ can be calculated through the measure matrix $\Phi \in \mathbf{R}^{n \times N}(n \ll N)$ as follows:

$$y = \Phi A_i \tag{1}$$

If the measured data is recovered through $\hat{A}_i = f(y) = \Psi y = \Psi \Phi A_i$, the estimated results can be obtained by optimization

$$min \parallel A_i - f(y) \parallel_{l_x} \tag{2}$$

$$s.t. \quad \Phi A_i = y \tag{3}$$

where $l_x(x = 0, 1, 2)$ are the norms of a vector, among which l_1 is usually used representing the number of non-zero coordinates. Then y can be easily measured and A can be reconstructed through y if the measure and reconstruct matrix (Φ and Ψ) are suitably designed. Therefore, all-to-all communications are not required which can help to reduce the measurement overhead. The topology graph can be estimated by Compress Sensing due to its large-scale and sparse feature. Furthermore, the propagation table and traffic loads can be measured in this way according to the measurement results in [12] and [2].

The topology that should be measured includes:

- Propagation table → Capacity matrix: the curve is quite certain for high frequency at near distance (do not consider PHY and MAC settings) [12], [13]
- PAs → LAs
- Traffic loads → Demands: the distribution of traffic loads has location differences across AAs and ToRs [2], [12]

Topology changes due to:

- Errors: nodes, links, miswiring
- wireless linking (links)
- Sleeping on-off (nodes)

B. Demand Response

according to the traffic measurement and analysis in VL2 and wireless flyways, traffic patterns and loads are dynamic and unpredictable. So demand response can be done by Lyapunov optimization for it only needs the information of current and previous state.

C. Capacity Scheduling

Multi-path provides extra selection of routing paths (just as over-provisioning)

D. Resource Allocation

high efficiency: good trade-off between reliability and capacity

III. MOTIVATION

With the increasing demand of traffic loads and user requirements, it is a great challenge to make data centers agile and energy-efficient. A basic solution is to allow dynamic resource allocation based on flexible data center networks. The recent development of 60GHz wireless technology opens a door for the deployment of flexible data centers. It provides a good choice of adding capacity to data centers with under-provisioning capacity design. However, this will lead to new problems such as topology control and capacity scheduling, apart from the wireless propagation and link adaption in wireless networks. The PHY and MAC standardization of 60GHz networks is still ongoing (WirelessHD and IEEE 802.11ad/WiGig), its application in data centers should be further explored for the specific feature of topology and services.

| (a) Provisioning for peak load | (b) Fixed underprovisioning | (c) Flexible underprovisioning |

Fig. 1. Capacity provisioning based on traffic loads

Generally, the capacity of data centers is designed for peak loads, as shown in Figure 1a, but the resources will be wasted during non-peak times [18]. On the other hand, if the capacity is designed in under-provisioning case as shown in Figure 1b, the peak requirements can not be satisfied, leading to potential revenue sacrifice or over-occupied errors. To address these problems, we can employ flexible capacity and dynamic demand response, i.e., 60GHz wireless linking to add capacity for peak loads and sleeping mechanism to improve performance/cost efficiency, as shown in Figure 1c.

The wireless-flexible data centers are generally composed of:

1) Topology and Addressing:
    a) Conflict Graph: Propagation table (wireless), capacity matrix (wired/wireless) and traffic loads table (wired/wireless)
    b) Topology Graph: Physical-specific ID Addresses (PAs), Logical-specific IP Addresses (LAs) and Application-specific IP Addresses (AAs)

i) PAs: providing distance information for 3D Beamforming
   ii) LAs: providing topology information for Capacity Adaption
   iii) AAs: providing traffic information for Demand Response
2) Routing and Scheduling:
   a) 3D Beamforming: adding additional capacity to neighbor ToRs according to PAs
   b) Demand Response: traffic estimation according to AAs
   c) Capacity Adaption: capacity scheduling according to LAs

IV. METHODOLOGY

The problems of topology control, addressing, routing, and scheduling also exit in traditional data centers, but it brings new challenges when wireless links are introduced. For instance, the topology graph is changing due to wireless links and on-demand response. Also, the capacity matrix and conflict graph of wireless links are closely related to the relative locations of servers, e.g., capacity of 6G for neighbor servers and 1/2G for non-neighbor servers. Furthermore, the measurement and mapping of topology graph, conflict graph and traffic loads is challenging for large scale data centers.
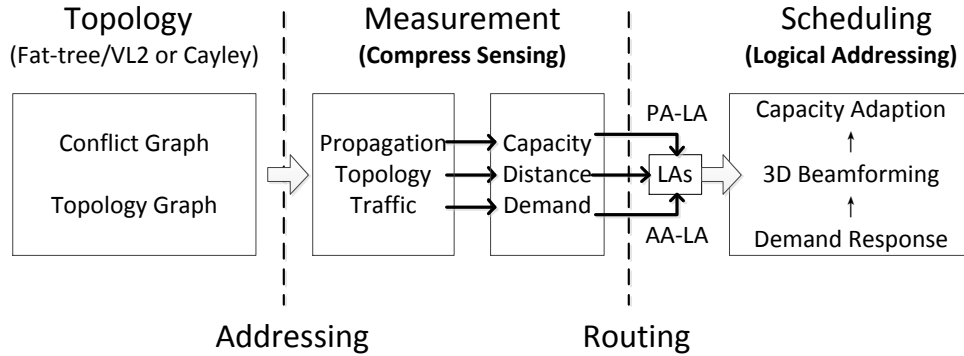

Fig. 2.   Framework of flexible data centers based on 60GHz wireless networks

Considering the scale of topology (conflict/topology graph and different addresses) and multi-layer scheduling (physical topology, capacity matrix and traffic patterns), we can employ **Compress Sensing** and **Logical Addressing** to address these problems.

Topology: multi-layer tree topology such as Fat-tree and VL2 (or cylindrical racks housing pie-shaped servers [14]) with ToRs.

Addressing: changing according to topology, PAs and AAs should be mapped to LAs to make 3D beamforming, demand response and capacity scheduling.

Scheduling: when the capacity matrix and traffic demands are mapped to LAs, we can determine the scheduling strategy including 3D beamforming and routing paths.

A. Compress Sensing

The conflict graph can be classified into the topology graph, since it represents the wireless connections which provide additional 1/2/6G capacity for different ToRs. It is challenging to measure such a large scale topology graph (PAs and AAs) which usually requires all-to-all communications. Since both conflict graph [12] and topology graph [15] are sparse, we can utilize Compress Sensing to reduce overhead, just as the famous example that using only 3 out of 10 words can reconstruct the whole account password.

In this way, we can use small number of samples to get the topology so that all-to-all communications are not necessary.

The requirements and reasons for Compress Sensing are:

1) topology graphs are **large**, and so comes the requirement to reduce measuring and mapping overhead
2) wireless propagation and traffic demand have **dynamic** furthers, so the measurement is challenging
3) topology graphs are **sparse**, which is the prior condition for compress sensing
4) topology graphs only have **local changes** (errors, on-off or wireless) in real-time operating

If the adjacency matrix of a given topology graph $G = (V, E)$ is $A^{N \times N}$, it can be represented by N vertex A_i. Then the estimated matrix $\hat{y} \in \mathbf{R}^n$ can be calculated through the measure matrix $\Phi \in \mathbf{R}^{n \times N}(n \ll N)$ as follows:

$$y = \Phi A_i \tag{4}$$

If the measured data is recovered through $\hat{A}_i = f(y) = \Psi y = \Psi \Phi A_i$, the estimated results can be obtained by optimization

$$min \parallel A_i - f(y) \parallel_{l_x} \tag{5}$$

$$s.t. \quad \Phi A_i = y \tag{6}$$

where $l_x(x = 0, 1, 2)$ are the norms of a vector, among which l_1 is usually used representing the number of non-zero coordinates. Then y can be easily measured and A can be reconstructed through y if the measure and reconstruct matrix (Φ and Ψ) are suitably designed. Therefore, all-to-all communications are not required which can help to reduce the measurement overhead. The topology graph can be estimated by Compress Sensing due to its large-scale and sparse feature. Furthermore, the propagation table and traffic loads can be measured in this way according to the measurement results in [12] and [2].

The topology that should be measured includes:

- Propagation table $\rightarrow$ Capacity matrix: the curve is quite certain for high frequency at near distance (do not consider PHY and MAC settings) [12], [13]
- PAs $\rightarrow$ LAs
- Traffic loads $\rightarrow$ Demands: the distribution of traffic loads has location differences across AAs and ToRs [2], [12]

Topology changes due to:

- Errors: nodes, links, miswiring
- wireless linking (links)
- Sleeping on-off (nodes)

B. Logical Addressing

Wireless links make data centers flexible, i.e., using topology control and capacity scheduling to make resource allocation responding to traffic patterns and requirements. So the problem is how to allocate wireless links to deal with the problem of demand response when topology graphs are changing. The capacity matrix is composed of wired and wireless links (direct links between neighbor nodes or 3D beamforming for remote nodes), which can be obtained by PAs:

- wired 10G
- wireless direct 6G
- wireless indirect 1/2G

Then the PAs (conflict and topology graphs) and demands estimation (AAs) can be mapped to LAs:

- PA to LA for 3D Beamforming
- AA to LA for Traffic estimation

When the topology graph and traffic demands are mapped to LAs, the capacity scheduling can be made responding to dynamic traffic loads. First, we can get the whole capacity matrix according to propagation table and topology graph. Second, greedy choice of flyways can be made according to capacity matrix an traffic demands. Finally, we can make wireless linking by direct or indirect beamforming if necessary, and make routing and scheduling further.

REFERENCES

[1] Mohammad Al-Fares, Alexander Loukissas, and Amin Vahdat. A scalable, commodity data center network architecture. ACM SIGCOMM '08, 2008.

[2] Albert Greenberg, James R. Hamilton, Navendu Jain, Srikanth Kandula, Changhoon Kim, Parantap Lahiri, David A. Maltz, Parveen Patel, and Sudipta Sengupta. VL2: a scalable and flexible data center network. ACM SIGCOMM '09, 2009.

[3] Chuanxiong Guo, Haitao Wu, Kun Tan, Lei Shi, Yongguang Zhang, and Songwu Lu. DCell: a scalable and fault-tolerant network structure for data centers. ACM SIGCOMM '08, 2008.

[4] Chuanxiong Guo, Guohan Lu, Dan Li, Haitao Wu, Xuan Zhang, Yunfeng Shi, Chen Tian, Yongguang Zhang, and Songwu Lu. BCube: a high performance, server-centric network architecture for modular data centers. ACM SIGCOMM '09, 2009.

[5] A. Singla, C.Y. Hong, L. Popa, and P.B. Godfrey. Jellyfish: Networking data centers randomly. USENIX NSDI '12, 2012.

[6] Jayaram Mudigonda, Praveen Yalagandula, Jeff Mogul, Bryan Stiekes, and Yanick Pouffary. NetLord: a scalable multi-tenant network architecture for virtualized datacenters. ACM SIGCOMM '11, 2011.

[7] Hitesh Ballani, Paolo Costa, Thomas Karagiannis, and Ant Rowstron. Towards predictable datacenter networks. ACM SIGCOMM '11, 2011.

[8] Lucian Popa, Arvind Krishnamurthy, Sylvia Ratnasamy, and Ion Stoica. FairCloud: sharing the network in cloud computing. ACM SIGCOMM '12, 2012.

[9] Guohui Wang, David G. Andersen, Michael Kaminsky, Konstantina Papagiannaki, T.S. Eugene Ng, Michael Kozuch, and Michael Ryan. c-Through: part-time optics in data centers. ACM SIGCOMM '10, 2010.

[10] Nathan Farrington, George Porter, Sivasankar Radhakrishnan, Hamid Hajabdolali Bazzaz, Vikram Subramanya, Yeshaiahu Fainman, George Papen, and Amin Vahdat. Helios: a hybrid electrical/optical switch architecture for modular data centers. ACM SIGCOMM '10, 2010.

[11] K. Chen, A. Singla, A. Singh, K. Ramachandran, L. Xu, Y. Zhang, X. Wen, and Y. Chen. OSA: An optical switching architecture for data center networks with unprecedented flexibility. USENIX NSDI '12, 2012.

[12] Daniel Halperin, Srikanth Kandula, Jitendra Padhye, Paramvir Bahl, and David Wetherall. Augmenting data center networks with multi-gigabit wireless links. ACM SIGCOMM '11, 2011.

[13] Xia Zhou, Zengbin Zhang, Yibo Zhu, Yubo Li, Saipriya Kumar, Amin Vahdat, Ben Y. Zhao, and Haitao Zheng. Mirror mirror on the ceiling: flexible wireless links for data centers. ACM SIGCOMM '12, 2012.

[14] Ji-Yong Shin, Emin Gün Sirer, Hakim Weatherspoon, and Darko Kirovski. On the feasibility of completely wireless datacenters. ACM ANCS '12, 2012.

[15] Kai Chen, Chuanxiong Guo, Haitao Wu, Jing Yuan, Zhenqian Feng, Yan Chen, Songwu Lu, and Wenfei Wu. Generic and automatic address configuration for data center networks. ACM SIGCOMM '10, 2010.

[16] Radhika Niranjan Mysore, Andreas Pamboris, Nathan Farrington, Nelson Huang, Pardis Miri, Sivasankar Radhakrishnan, Vikram Subramanya, and Amin Vahdat. PortLand: a scalable fault-tolerant layer 2 data center network fabric. ACM SIGCOMM '09, 2009.

[17] Xin Wu, Daniel Turner, Chao-Chih Chen, David A. Maltz, Xiaowei Yang, Lihua Yuan, and Ming Zhang. NetPilot: automating datacenter network failure mitigation. ACM SIGCOMM '12, 2012.

[18] Michael Armbrust, Armando Fox, Rean Griffith, Anthony D. Joseph, Randy Katz, Andy Konwinski, Gunho Lee, David Patterson, Ariel Rabkin, Ion Stoica, and Matei Zaharia. A view of cloud computing. *Commun. ACM*, 53(4):50–58, April 2010.